

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
df = pd.read_csv("/content/diabetes_prediction_dataset.csv")
```

```
df.head()
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	Female	80.0	0	1	never	25.19	6.6	140	0
1	Female	54.0	0	0	No Info	27.32	6.6	80	0
2	Male	28.0	0	0	never	27.32	5.7	158	0
3	Female	36.0	0	0	current	23.45	5.0	155	0
4	Male	76.0	1	1	current	20.14	4.8	155	0

```
df.shape
```

(100000, 9)

```
df.describe()
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level	diabetes
count	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000
mean	41.885856	0.07485	0.039420	27.320767	5.527507	138.058060	0.085000
std	22.516840	0.26315	0.194593	6.636783	1.070672	40.708136	0.278883
min	0.080000	0.00000	0.000000	10.010000	3.500000	80.000000	0.000000
25%	24.000000	0.00000	0.000000	23.630000	4.800000	100.000000	0.000000
50%	43.000000	0.00000	0.000000	27.320000	5.800000	140.000000	0.000000
75%	60.000000	0.00000	0.000000	29.580000	6.200000	159.000000	0.000000
max	80.000000	1.00000	1.000000	95.690000	9.000000	300.000000	1.000000

```
df['diabetes'].value_counts()
```

0 91500
1 8500
Name: diabetes, dtype: int64

```
df['gender'].value_counts()
```

Female 58552
Male 41430
Other 18
Name: gender, dtype: int64

```
#Converting categorical variable to numeric
```

```
import pandas as pd
gender_mapping = {'Male': 1, 'Female': 0, 'Other': 2}
```

```
df['numeric_gender'] = df['gender'].map(gender_mapping)
print(df)
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	\
0	Female	80.0	0	1	never	25.19	
1	Female	54.0	0	0	No Info	27.32	
2	Male	28.0	0	0	never	27.32	
3	Female	36.0	0	0	current	23.45	
4	Male	76.0	1	1	current	20.14	
...	
99995	Female	80.0	0	0	No Info	27.32	
99996	Female	2.0	0	0	No Info	17.37	
99997	Male	66.0	0	0	former	27.83	

```

99998 Female 24.0      0      0      never 35.42
99999 Female 57.0      0      0      current 22.43

```

```

      HbA1c_level  blood_glucose_level  diabetes  numeric_gender
0           6.6           140           0           0
1           6.6           80           0           0
2           5.7          158           0           1
3           5.0          155           0           0
4           4.8          155           0           1
...         ...         ...         ...         ...
99995         6.2           90           0           0
99996         6.5          100           0           0
99997         5.7          155           0           1
99998         4.0          100           0           0
99999         6.6           90           0           0

```

```
[100000 rows x 10 columns]
```

```
df.head()
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood
0	Female	80.0	0	1	never	25.19	6.6	
1	Female	54.0	0	0	No Info	27.32	6.6	
2	Male	28.0	0	0	never	27.32	5.7	
3	Female	36.0	0	0	current	23.45	5.0	
4	Male	76.0	1	1	current	20.14	4.8	

```
df['numeric_gender'].value_counts()
```

```

0    58552
1    41430
2         18
Name: numeric_gender, dtype: int64

```

```
df['smoking_history'].value_counts()
```

```

No Info    35816
never      35095
former     9352
current    9286
not current 6447
ever       4004
Name: smoking_history, dtype: int64

```

```
#Converting categorical variable to numeric
```

```
import pandas as pd
```

```
smoking_history_numeric = {'No Info': 0, 'never': 1, 'former': 2, 'current': 3, 'not current':4, 'ever':5}
```

```
df['smoking_history_numeric'] = df['smoking_history'].map(smoking_history_numeric)
print(df)
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	\
0	Female	80.0	0	1	never	25.19	
1	Female	54.0	0	0	No Info	27.32	
2	Male	28.0	0	0	never	27.32	
3	Female	36.0	0	0	current	23.45	
4	Male	76.0	1	1	current	20.14	
...	
99995	Female	80.0	0	0	No Info	27.32	
99996	Female	2.0	0	0	No Info	17.37	
99997	Male	66.0	0	0	former	27.83	
99998	Female	24.0	0	0	never	35.42	
99999	Female	57.0	0	0	current	22.43	

	HbA1c_level	blood_glucose_level	diabetes	numeric_gender	\
0	6.6	140	0	0	
1	6.6	80	0	0	
2	5.7	158	0	1	
3	5.0	155	0	0	
4	4.8	155	0	1	
...	
99995	6.2	90	0	0	
99996	6.5	100	0	0	
99997	5.7	155	0	1	

```

99998      4.0      100      0      0
99999      6.6      90      0      0

```

```

smoking_history_numeric
0      1
1      0
2      1
3      3
4      3
...    ...
99995      0
99996      0
99997      2
99998      1
99999      3

```

```
[100000 rows x 11 columns]
```

```
df.head()
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood
0	Female	80.0	0	1	never	25.19	6.6	
1	Female	54.0	0	0	No Info	27.32	6.6	
2	Male	28.0	0	0	never	27.32	5.7	
3	Female	36.0	0	0	current	23.45	5.0	
4	Male	76.0	1	1	current	20.14	4.8	

```

# Drop the specified columns
columns_to_remove = ['gender', 'smoking_history']
df = df.drop(columns=columns_to_remove)

```

```
df.head()
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level	diabetes
0	80.0	0	1	25.19	6.6	140	0
1	54.0	0	0	27.32	6.6	80	0
2	28.0	0	0	27.32	5.7	158	0
3	36.0	0	0	23.45	5.0	155	0
4	76.0	1	1	20.14	4.8	155	0

```

# separating the data and labels
X = df.drop(columns = 'diabetes', axis=1)
Y = df['diabetes']

```

Data Standardization

```

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
standardized_data = scaler.fit_transform(X)

```

```
print(standardized_data)
```

```

[[ 1.69270354 -0.28443945  4.93637859 ...  0.04770422 -0.84104674
 -0.19575171]
 [ 0.53800643 -0.28443945 -0.20257766 ... -1.42620999 -0.84104674
 -0.90848322]
 [-0.61669069 -0.28443945 -0.20257766 ...  0.48987848  1.18723364
 -0.19575171]
 ...
 [ 1.07094356 -0.28443945 -0.20257766 ...  0.41618277  1.18723364
  0.5169798 ]
 [-0.7943364 -0.28443945 -0.20257766 ... -0.93490525 -0.84104674
 -0.19575171]

```

```
[ 0.67124071 -0.28443945 -0.20257766 ... -1.18055762 -0.84104674
 1.22971132]]
```

```
X = standardized_data
Y= df['diabetes']
```

Train Test Split

```
from sklearn.model_selection import train_test_split

X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size = 0.2, stratify=Y, random_state=2)

print(X.shape, X_train.shape, X_test.shape)

(100000, 8) (80000, 8) (20000, 8)
```

Training the Model

```
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report

# Initialize SVM classifier
svm_classifier = SVC(kernel='linear', C=1.0, random_state=42)

svm_classifier.fit(X_train, Y_train)
```

▼
SVC

SVC(kernel='linear', random_state=42)

```
# Make predictions on the testing set
y_pred = svm_classifier.predict(X_test)

# accuracy score on the training data
X_train_prediction = svm_classifier.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)

print('Accuracy score of the training data : ', training_data_accuracy)

Accuracy score of the training data :  0.96045

# accuracy score on the testing data
X_test_prediction = svm_classifier.predict(X_test)
testing_data_accuracy = accuracy_score(X_test_prediction, Y_test)

print('Accuracy score of the testing data : ', testing_data_accuracy)

Accuracy score of the testing data :  0.96165
```

```
# Evaluate the performance of the classifier
accuracy = accuracy_score(Y_test, y_pred)
classification_report_output = classification_report(Y_test, y_pred)

# Print the results
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:")
print(classification_report_output)
```

```
Accuracy: 0.96
Classification Report:
              precision    recall  f1-score   support

     0       0.96       1.00       0.98       18300
     1       0.92       0.60       0.73       1700

 accuracy         0.96         0.96         0.96       20000
  macro avg       0.94         0.80         0.85       20000
 weighted avg     0.96         0.96         0.96       20000
```

Making a predictive System

```
df.head(10)
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level	diabetes
0	80.0	0	1	25.19	6.6	140	0
1	54.0	0	0	27.32	6.6	80	0
2	28.0	0	0	27.32	5.7	158	0
3	36.0	0	0	23.45	5.0	155	0
4	76.0	1	1	20.14	4.8	155	0
5	20.0	0	0	27.32	6.6	85	0
6	44.0	0	0	19.31	6.5	200	1
7	79.0	0	0	23.86	5.7	85	0
8	42.0	0	0	33.64	4.8	145	0
9	32.0	0	0	27.32	5.0	100	0

```
input_data = (44.0,1,1,19.31,6.5,200,1,1)
```

```
# changing the input_data to numpy array
input_data_as_numpy_array = np.asarray(input_data)
```

```
# reshape the array as we are predicting for one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
```

```
# standardize the input data
std_data = scaler.transform(input_data_reshaped)
print(std_data)
```

```
[[ 0.09389215  3.51568677  4.93637859 -1.20703158  0.90830598  1.52161842
  1.18723364 -0.19575171]]
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was f
warnings.warn(
```

```
prediction = svm_classifier.predict(std_data)
print(prediction)
```

```
[1]
```

```

if (prediction[0] == 0):
    print('The person is not diabetic')
else:
    print('The person is diabetic')

```

Saving the trained model

```

import pickle

filename = 'trained_model.sav'

pickle.dump(svm_classifier, open(filename, 'wb'))

```

Loading the saved model

```

loaded_model = pickle.load(open('trained_model.sav', 'rb'))

```

```

input_data = (44.0,1,1,19.31,6.5,200,1,1)

# changing the input_data to numpy array
input_data_as_numpy_array = np.asarray(input_data)

# reshape the array as we are predicting for one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = svm_classifier.predict(input_data_reshaped)
print(prediction)

if (prediction[0] == 0):
    print('The person is not diabetic')
else:
    print('The person is diabetic')

[1]
The person is diabetic

```